

CSCI 4210 — Operating Systems
CSCI 6140 — Computer Operating Systems
Project 2 (document version 1.5)
Memory Management and Virtual Memory

Overview

- This project is due by 11:59:59 PM on Monday, December 12, 2016.
- This project will count as 12% of your final course grade.
- This project is to be completed **individually** or in a team of at most three students. Do not share your code with anyone else.
- You **must** use one of the following programming languages: C, C++, Java, or Python. If using Java, name your main Java file `Project2.java`. (**v1.5**) If using Python, name your main Python file `project2.py`.
- Your program **must** successfully compile and run on Ubuntu v14.04.5 LTS or higher.
- If you use C or C++, your program **must** successfully compile via `gcc` or `g++` with absolutely no warning messages when the `-Wall` (i.e., warn all) compiler option is used. We will also use `-Werror`, which will treat all warnings as critical errors.
- For Java and Python, be sure no warning messages occur during compilation and/or interpretation.

Project Specifications

In this second project, you will simulate a variety of memory management systems, including contiguous, non-contiguous, and virtual memory. In an operating system, each process has specific memory requirements. These memory requirements are met (or not) based on whether free memory is available to fulfill such requirements.

Representing Physical (Main) Memory

For contiguous and non-contiguous memory schemes, to represent physical memory, use a data structure of your choice to represent a memory that contains a configurable number of equally sized “memory units” or frames.

As an example, memory may be represented as shown below (also see examples in the in-class notes). When you display your simulated memory, use the format below, showing exactly 32 frames per line. Also, as a default, simulate a memory consisting of 256 frames (i.e., eight output lines). These two values should be configurable and easily tunable.

```

=====
AAAAAAAABBBBBBBBBBBBBBBBBBBBBBBB
BBBBBBBBBBBBBBB.....
.....DDDDDDDD....
.....
.....
.....HHHHHHHHHHHHHHHHHHHHHHHH
HHH.....
.....
=====

```

More specifically, use the following character mappings:

- A '.' character represents a free memory frame.
- An ASCII character in the range 'A'-'Z' (i.e., uppercase letters only) represents a user process memory frame. Note that we will ignore operating system process memory frames.

Input File

For contiguous and non-contiguous memory schemes, details of user processes must be read from an input file, which is to be specified as the first command-line argument of your simulation. The input file is formatted as shown below. (**v1.1**) Any line beginning with a # character is ignored (these lines are comments). Further, all blank lines are also ignored, including lines containing only whitespace characters.

```

N
proc1 p1_mem p1_arr_time_1/p1_run_time_1 ... p1_arr_time_a/p1_run_time_a
proc2 p2_mem p2_arr_time_1/p2_run_time_1 ... p1_arr_time_b/p1_run_time_b
...
procN pN_mem pN_arr_time_1/pN_run_time_1 ... p1_arr_time_z/p1_run_time_z

```

Here, the overall number of processes to simulate (N) is specified first on a line by itself (note that this allows for dynamic memory allocation to occur). Each **proc#** value specifies a single character in the range 'A'-'Z' that will uniquely identify the given process. Each **p#_mem** value specifies the required (fixed) number of memory frames.

Each **p#_arr_time_?/p#_run_time_?** pair specifies a corresponding pair of arrival and run times for the given process. You may assume that each process has an increasing set of non-zero arrival times. You may also assume that each run time is greater than zero. In other words, you do not need to validate these rules in your code.

Further, assume that the maximum number of processes in the simulation will be 26. Do not assume that processes will be given in alphabetical order.

Below is an example input file (note that values are delimited by one or more space or TAB characters):

```
20
A 45 0/350 400/50
B 28 0/2650
C 58 0/950 1100/100
D 86 0/650 1350/450
E 14 0/1400
F 24 100/380 500/475
G 13 435/815
J 46 550/900
etc.
```

Note that for the contiguous memory management scheme, when defragmentation occurs, these numbers must be automatically adjusted by extending all future arrival times accordingly. As an example, if defragmentation takes 300 time units, then all pending arrival times should increase by 300.

Also note that any “ties” that occur should be handled using the alphabetical order of process IDs.

Contiguous Memory Management

For this portion of the simulation, each process's memory requirements must be met by the various contiguous memory allocation schemes that we have studied. The algorithms you must simulate are the *next-fit*, *best-fit*, and *worst-fit* algorithms, each of which is described in more detail below.

Note that we will only use a dynamic partitioning scheme here, meaning that your data structure needs to maintain a list containing (1) where each process is allocated, (2) how much contiguous memory each process uses, and (3) where and how much free memory is available (i.e., where each free partition is).

As an example, consider the simulated memory shown below.

```
=====
AAAAAAAABBBBBBBBBBBBBBBBBBBBBB
BBBBBBBBBBBBBBB.....
.....DDDDDDDD....
.....
.....
.....HHHHHHHHHHHHHHHHHHHHHH
HHH.....
.....
=====
```

In the above example diagram, four processes are allocated in four dynamically allocated partitions. Further, there are three free partitions (i.e., between processes B and D, between processes D and H, and between process H and the “bottom” of memory).

Placement Algorithms

When a process Q wishes to enter the system, memory must first be allocated dynamically. More specifically, a dynamic partition must be created out of an existing free partition.

For the next-fit algorithm, process Q is placed in the first free partition available found by scanning from the end of the most recently placed process.

For the best-fit algorithm, process Q is placed in the smallest free partition available in which process Q fits. If a “tie” occurs, use the free partition closer to the “top” of memory.

For the worst-fit algorithm, process Q is placed in the largest free partition available in which process Q fits. If a “tie” occurs, use the free partition closer to the “top” of memory.

For all of these placement algorithms, the memory scan covers the entire memory. If necessary (i.e., if the scan hits the “bottom” of memory), the scan continues from the “top” of memory.

If no suitable free partition is available, then an out-of-memory error occurs, at which point defragmentation is required.

Defragmentation

If a process is unable to be placed into memory, defragmentation occurs if the overall total of free memory is sufficient to fit the given process. In such cases, processes are relocated as necessary, starting from the top of memory. If there is not enough memory to admit the given process, your simulation must skip defragmentation and deny the request and move on to the next process.

Note that the given process is only skipped for the given requested interval. It may be admitted to the system at a later interval. For example, given process Q below, if the process is unable to be added in interval 200–580, it might still be successfully added in interval 720–975.

Q 47 200/380 720/255

Once defragmentation has started, it will run through to completion, at which point, the process that triggered defragmentation will be admitted to the system.

While defragmentation is running, no other processes may be executing (i.e., all processes are essentially in a suspended state) until all relocations are complete. Thus, arrivals and exits are suspended during defragmentation.

Note that the time to move **one frame of memory** is defined as `t_memmove` and is measured in milliseconds. Assume that the default value of `t_memmove` is 1.

Given the memory shown previously, the results of defragmentation will be as follows (i.e., processes D and H are moved upwards in memory):

```
=====
AAAAAAAABBBBBBBBBBBBBBBBBBBBBBBB
BBBBBBBBBBBBBBBBDDDDDDDDHHHHHHHH
HHHHHHHHHHHHHHHHHHHH.....
.....
.....
.....
.....
.....
=====
```

Simulation and Output Requirements

For the given input file, simulate each of the three placement algorithms, displaying output for each “interesting” event that occurs. Interesting events are:

- Simulation start
- Process arrival
- Process placement in physical memory
- Process exit from physical memory
- Start of defragmentation
- End of defragmentation (i.e., process(es) moved)
- Simulation end

As with previous projects, your simulator must keep track of elapsed time t (measured in milliseconds), which is initially set to zero. As your simulation proceeds, based on the input file, t advances to each “interesting” event that occurs, displaying a specific line of output describing each event.

Note that your simulator output should be entirely deterministic. To achieve this, your simulator must output each “interesting” event that occurs using the format shown below.

```
time <t>ms: <event-details>
```

When memory changes, you should display the full simulated memory.

Given the example input file at the bottom of page 2 and default configuration parameters, your simulator output would be as follows:

```
time 0ms: Simulator started (Contiguous -- Next-Fit)
time 0ms: Process A arrived (requires 45 frames)
time 0ms: Placed process A:
=====
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAA.....
.....
.....
.....
.....
.....
.....
.....
=====
```

time 0ms: Process B arrived (requires 28 frames)
time 0ms: Placed process B:

```
=====
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAABBBBBBBBBBBBBBBB
BBBBBBBBB.....
.....
.....
.....
.....
=====
```

time 0ms: Process C arrived (requires 58 frames)
time 0ms: Placed process C:

```
=====
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAABBBBBBBBBBBBBBBB
BBBBBBBBBCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCCCCC
CCC.....
.....
.....
=====
```

time 0ms: Process D arrived (requires 86 frames)
time 0ms: Placed process D:

```
=====
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAABBBBBBBBBBBBBBBB
BBBBBBBBBCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCCCCC
CCDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDD.....
.....
=====
```

time 0ms: Process E arrived (requires 14 frames)
time 0ms: Placed process E:

```
=====
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAABBBBBBBBBBBBBBBB
BBBBBBBBBCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCCCCC
CCDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDEEEEEEE
EEEEEE.....
=====
```

time 100ms: Process F arrived (requires 24 frames)
time 100ms: Placed process F:

```
=====
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAABBBBBBBBBBBBBBBB
BBBBBBBBBCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCCCCCC
CCDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDEEEEEEE
EEEEEEFFFFFFFFFFFFFFFFFFFFFFF.
=====
```

time 350ms: Process A removed:

```
=====
.....
.....BBBBBBBBBBBBBBBBBBBBB
BBBBBBBBBCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCCCCCC
CCDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDEEEEEEE
EEEEEEFFFFFFFFFFFFFFFFFFFFFFF.
=====
```

time 400ms: Process A arrived (requires 45 frames)

time 400ms: Placed process A:

```
=====
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAABBBBBBBBBBBBBBBB
BBBBBBBBBCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCCCCCC
CCDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDEEEEEEE
EEEEEEFFFFFFFFFFFFFFFFFFFFFFF.
=====
```

time 435ms: Process G arrived (requires 13 frames)

time 435ms: Cannot place process G -- skipped!

```
=====
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAABBBBBBBBBBBBBBBB
BBBBBBBBBCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCCCCCC
CCDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDEEEEEEE
EEEEEEFFFFFFFFFFFFFFFFFFFFFFF.
=====
```



```

time 450ms: Process A removed:
=====
.....
.....BBBBBBBBBBBBBBBBBB
BBBBBBBBBBBBBBBBBBBBBBBB
CCCCCCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCC
CCDDDDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDEEEEEEE
EEEEEEFFFFFFFFFFFFFFFFFFFFFFFFF.
=====

time 480ms: Process F removed:
=====
.....
.....BBBBBBBBBBBBBBBBBB
BBBBBBBBBBBBBBBBBBBBBBBB
CCCCCCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCC
CCDDDDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDEEEEEEE
EEEEEE.....
=====

time 500ms: Process F arrived (requires 24 frames)
time 500ms: Placed process F:
=====
.....
.....BBBBBBBBBBBBBBBBBB
BBBBBBBBBBBBBBBBBBBBBBBB
CCCCCCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCC
CCDDDDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDEEEEEEE
EEEEEEFFFFFFFFFFFFFFFFFFFFFFFFF.
=====

time 550ms: Process J arrived (requires 46 frames)
time 550ms: Cannot place process J -- starting defragmentation
time 760ms: Defragmentation complete (moved 210 frames: B, C, D, E, F)
=====
BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBB
CCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDEEEEEEEEEEEEEEEFFFFFFF
FFFFFFFFFFFFFFFFFFFF.....
.....
=====

```

```

time 760ms: Placed process J:
=====
BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBCCCC
CCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDEEEEEEEEEEEEEEEEEFFFFF
FFFFFFFFFFFFFFFFJJJJJJJJJJJJJJJJJJJJJJJJJJ
JJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJ
=====

time 860ms: Process D removed:
=====
BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBCCCC
CCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCC.....
.....
.....
.....EEEEEEEEEEEEEEEEFFFFF
FFFFFFFFFFFFFFFFJJJJJJJJJJJJJJJJJJJJJJJJJJ
JJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJ
=====

...
time ###ms: Simulator ended (Contiguous -- Next-Fit)

time 0ms: Simulator started (Contiguous -- Best-Fit)
time 0ms: Process A arrived (requires 45 frames)
time 0ms: Placed process A:
=====
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAA.....
.....
.....
.....
.....
.....
.....
=====

...
time ###ms: Simulator ended (Contiguous -- Best-Fit)

```

```

time 0ms: Simulator started (Contiguous -- Worst-Fit)
time 0ms: Process A arrived (requires 45 frames)
time 0ms: Placed process A:
=====
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAA.....
.....
.....
.....
.....
.....
.....
=====
...
time ###ms: Simulator ended (Contiguous -- Worst-Fit)

```

Non-contiguous Memory Management

Extend the above contiguous memory management simulation by next simulating a non-contiguous memory management scheme in which you use a page table to map each logical page of a process to a physical frame of memory.

The key difference here is that defragmentation is no longer necessary. To place pages into frames of physical memory, use a simple first-fit approach, as shown in the example below.

```
time 0ms: Simulator started (Non-contiguous)
time 0ms: Process A arrived (requires 45 frames)
time 0ms: Placed process A:
=====
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAA.....
.....
.....
.....
.....
.....
.....
=====
...
time 500ms: Process F arrived (requires 24 frames)
time 500ms: Placed process F:
=====
FFFFFFFFFFFFFFFFFFFFFFFF.....
.....BBBBBBBBBBBBBBBBBBBB
BBBBBBBBBBBBBBBBBBBBBBBBBBBB
CCCCCCCCCCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCCCCCC
CCDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDEEEEEEE
EEEEEE.....
=====
time 550ms: Process J arrived (requires 46 frames)
time 550ms: Placed process J:
=====
FFFFFFFFFFFFFFFFFFFFFFFFJJJJJJJJ
JJJJJJJJJJJJBBBBBBBBBBBBBBBBBB
BBBBBBBBBBBBBBBBBBBBBBBBBBBB
CCCCCCCCCCCCCCCCCCCCCCCCCCCC
CCCCCCCCCCCCCCCCCCCCCCCCCCCC
CCDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDDDDDD
DDDDDDDDDDDDDDDDDDDDDDDEEEEEEE
EEEEEEJJJJJJJJJJJJJJJJJJJJJJ
=====
...
time ###ms: Simulator ended (Non-contiguous)
```

Virtual Memory Management

For this part of the simulation, you will focus on page references strings and page replacement algorithms. The second command-line argument to your simulation specifies a file that contains a page reference string for a given process. Open and read the specified file for this part of your simulation. The file contains a page reference string, using spaces to delimit references. A snippet of the sample file given on the course website is shown below. Also note that there is simple Java code on the course website for generating random page reference strings.

```
7 3 3 3 3 3 3 3 1 1 8 3 3 3 3 3 3 3 6 6 6 6 6 6 6 7 7 7 7 7 3 3 3 3 3 3 3 3 8
8 8 8 8 8 8 5 5 5 5 2 2 2 2 2 2 2 2 2 7 7 4 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
3 3 3 3 3 3 3 3 5 5 5 5 5 5 5 5 5 6 6 6 6 6 5 5 5 1 2 2 5 5 5 5 7 7 7 7 3 3 3 3
...
```

In your simulation, the number of frames is initially set to F , which defaults to 3.

For the given page reference string (which we assume is associated with a single process), simulate the OPT, LRU, and LFU algorithms, each of which is described below.

The Optimal (OPT) algorithm is a forward-looking algorithm that selects the “victim” page by identifying the frame that will be accessed the longest time in the future (or not at all). If multiple pages are identified, this “tie” is broken by selecting the lowest-numbered page.

The Least-Recently Used (LRU) algorithm is a backward-looking algorithm that selects the “victim” page by identifying the frame that has the oldest access time.

The Least-Frequently Used (LFU) algorithm is a backward-looking algorithm that selects the “victim” page by identifying the frame with the lowest number of accesses. When a page fault occurs for a given page, its reference count is set (or reset) to 1; each subsequent access increments this reference count. If multiple pages are identified, this “tie” is broken by selecting the lowest-numbered page.

Simulation and Output Requirements

Append to the previous output your virtual memory results. In your output, show the fixed N -frame memory after each page fault; indicate page faults as shown below.

```
Simulating OPT with fixed frame size of 3
referencing page 7 [mem: 7 . .] PAGE FAULT (no victim page)
referencing page 3 [mem: 7 3 .] PAGE FAULT (no victim page)
referencing page 1 [mem: 7 3 1] PAGE FAULT (no victim page)
referencing page 8 [mem: 7 3 8] PAGE FAULT (victim page 1)
referencing page 6 [mem: 7 3 6] PAGE FAULT (victim page 8)
...
End of OPT simulation (### page faults)
```

```
Simulating LRU with fixed frame size of 3
```

```
referencing page 7 [mem: 7 . .] PAGE FAULT (no victim page)
...
End of LRU simulation (### page faults)

Simulating LFU with fixed frame size of 3
referencing page 7 [mem: 7 . .] PAGE FAULT (no victim page)
...
End of LFU simulation (### page faults)
```

Submission Instructions

To submit your assignment (and also perform final testing of your code), please use Submittity, the homework submission server. The URL is on the course website.

If you are submitting a team project, please have each team member submit the same submission (just to be sure everyone gets a grade). Also be sure to include all names and RCS IDs in comments at the top of each source file.